

**Homework 1 Solutions***Please do not distribute.***1. Bias-variance decomposition (15 points)**

Prove the claim in class:

$$\begin{aligned}\mathbb{E}_D \mathbb{E}_{X,Y}[(Y - \hat{h}_D(X))^2] &= \mathbb{E}_X[(\mathbb{E}_D[\hat{h}_D(X)] - h^*(X))^2] \\ &\quad + \mathbb{E}_X \mathbb{E}_D[(\hat{h}_D(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])^2] \\ &\quad + \mathbb{E}_{X,Y}[(Y - h^*(X))^2].\end{aligned}$$

**Solution:** Split  $Y - \hat{h}_D(X)$  to  $Y - h^*(X) + h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)] + \mathbb{E}_{D'}[\hat{h}'_D(X)] - \hat{h}_D(X)$ , we have

$$\begin{aligned}\mathbb{E}_D \mathbb{E}_{X,Y}[(Y - \hat{h}_D(X))^2] &= \mathbb{E}_{X,Y}[(\mathbb{E}_D[\hat{h}_D(X)] - h^*(X))^2] \\ &\quad + \mathbb{E}_{X,Y} \mathbb{E}_D[(\hat{h}_D(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])^2] \\ &\quad + \mathbb{E}_{X,Y} \mathbb{E}_D[(Y - h^*(X))^2] \\ &\quad + 2\mathbb{E}_{X,Y} \mathbb{E}_D[(Y - h^*(X))(h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])] \\ &\quad + 2\mathbb{E}_{X,Y} \mathbb{E}_D[(Y - h^*(X))(\hat{h}_D(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])] \\ &\quad + 2\mathbb{E}_{X,Y} \mathbb{E}_D[(\hat{h}_D(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])(h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])]\end{aligned}$$

The first three terms simplifies to the three terms in the desired equation. Now it suffices to prove the last three terms all equal zero.

$$\begin{aligned}&\mathbb{E}_{X,Y} \mathbb{E}_D [(Y - h^*(X))(h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])] \\ &= \mathbb{E}_X \mathbb{E}_{Y|X} [\mathbb{E}_D [(Y - h^*(X))(h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)])] | X] \\ &= \mathbb{E}_X \mathbb{E}_{Y|X} [(Y - h^*(X))(h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)]) | X]\end{aligned}$$

When  $X$  is fixed,  $h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)]$  is constant, so

$$\begin{aligned}&\mathbb{E}_X \mathbb{E}_{Y|X} [(Y - h^*(X))(h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)]) | X] \\ &= (h^*(X) - \mathbb{E}_{D'}[\hat{h}'_D(X)]) \mathbb{E}_{Y|X} [Y - h^*(X) | X]\end{aligned}$$

Since  $h^*(X) = \mathbb{E}_{Y|X} [Y | X]$ , the last term is zero. The other two terms can be shown to be zero in a similar manner.

**2. MLE (15 points)** Let  $\mathbf{x}$  be a sample from some distribution with parameter  $\theta$ . Let  $L(\theta|\mathbf{x})$  be the likelihood function for  $\theta \in \Theta$ . Let  $g : \Theta \rightarrow \mathbb{R}$  be a function and  $\tilde{L}(\xi|\mathbf{x}) := \sup_{\theta: g(\theta)=\xi} L(\theta|\mathbf{x})$  be the induced likelihood function. Suppose  $\hat{\theta}$  is the MLE of  $\theta$ , prove that  $g(\hat{\theta})$  is the MLE for  $g(\theta)$ .

**Solution:** Let  $\hat{\xi}$  be the MLE for  $g(\theta)$ . Then

$$\begin{aligned}\tilde{L}(\hat{\xi}|\mathbf{x}) &= \sup_{\xi} \sup_{\theta: g(\theta)=\xi} L(\theta|\mathbf{x}) \\ &= \sup_{\theta} L(\theta|\mathbf{x}) \\ &= L(\hat{\theta}|\mathbf{x}) \\ &= \sup_{g(\theta)=g(\hat{\theta})} L(\theta|\mathbf{x}) \\ &= \tilde{L}(g(\hat{\theta})|\mathbf{x})\end{aligned}$$

So  $g(\hat{\theta})$  is the MLE for  $g(\theta)$ .

### 3. Maximum likelihood estimation (30 points)

Suppose we have  $n$  i.i.d. samples  $X_1, \dots, X_n$  from the following distributions. Find the MLE of the required parameters.

a.(5 points) Unif(0,  $\theta$ ). Find MLE for  $\theta$ .

b.(5 points)  $N(\mu, \sigma^2)$  with both parameters unknown. Find MLE for  $\sigma^2$ .

c.(10 points)  $N(0, \Sigma)$ . Find the MLE for  $\Sigma$ .

d.(10 points)  $N(0, \Sigma)$ . Find the MLE for  $\Theta = \Sigma^{-1}$ . Please also provide conditions that such the MLE exists.

**Solution:** (a)  $\max_{1 \leq i \leq n} X_i$ .

(b)  $\frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$  where  $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ .

(c)  $\frac{1}{n} \sum_{i=1}^n X_i X_i^\top$

(d)  $(\frac{1}{n} \sum_{i=1}^n X_i X_i^\top)^{-1}$  given  $\frac{1}{n} \sum_{i=1}^n X_i X_i^\top$  is invertible.

### 4. Convex optimization (10 points)

Suppose that  $f: \mathbb{R}^d \mapsto \mathbb{R}$  is a convex and differentiable function. Show that  $x$  is a minimizer of  $f$  if and only if  $\nabla f(x) = 0$ .

**Solution:** Suppose  $\nabla f(x) = 0$ , then for any  $y \in \mathbb{R}^d$ ,  $f(y) \geq f(x) + \nabla f(x)^\top (y - x) = f(x)$  because of convexity.

Suppose  $x$  is a minimizer, take any unit vector  $v$ . From the definition of gradient we know

$$\lim_{t \rightarrow 0} \frac{f(x + tv) - f(x)}{t} = \nabla f(x)^\top v$$

Since the limit exists and  $\frac{f(x+tv)-f(x)}{t} > 0$  for  $t > 0$ ,  $\frac{f(x+tv)-f(x)}{t} < 0$  for  $t < 0$ , the limit has to be zero.  $\nabla f(x)^\top v = 0$  for all unit vector  $v$  implies  $\nabla f(x) = 0$ .

### 5. Programming assignment: Empirical risk minimization (30 points)

Given a sample data set

$D$  and loss function  $\ell(y, \hat{y})$ , the *empirical risk* (or empirical loss) of a hypothesis  $h$  is defined as the sample mean of the loss on  $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$ :

$$\hat{R}_D(h) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, h(x_i))$$

where  $h(x_i)$  is the predicted label for  $x_i$ . In *empirical risk minimization*, we use the empirical risk  $\hat{R}_D(h)$  as an estimator of the minimal true risk, a.k.a. the *Bayes risk*, defined as

$$R^* := \min_h \mathbb{E}_{X,Y} \ell(Y, h(X)).$$

Given hypothesis class  $\mathcal{H}$  (e.g. a collection of predictors), denote the empirical risk estimator as

$$\hat{h}_D := \operatorname{argmin}_{h \in \mathcal{H}} \hat{R}_D(h).$$

The expected error of the empirical risk estimator  $\hat{h}_D$  is

$$\mathbb{E}_D \hat{R}_D(\hat{h}_D) - R^* = \underbrace{\mathbb{E}_D \hat{R}_D(\hat{h}_D) - \min_{h \in \mathcal{H}} R(h)}_{\text{estimation error}} + \underbrace{\min_{h \in \mathcal{H}} R(h) - R^*}_{\text{approximation error}}$$

Hereby, the *estimation error* is due to error caused by  $n$  training samples instead of full knowledge of the joint distribution of  $X, Y$ , and the *approximation error* is due to restricting our attention to model class  $\mathcal{H}$ .

In this question, we will investigate the trade-off between *estimation error* and *approximation error* given different collections of predictors.

Use the following code to generate a dataset:

```
#python
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error

def data_generator(n_samples):
    x = np.random.uniform(-10, 10, n_samples)
    y = np.cos(0.5 + np.exp(-x)) + 1/(1 + np.exp(-x))
    noise = np.random.normal(0, 0.01, n_samples)
    y += noise
    return x, y


complete_X, complete_Y = data_generator(5000)
train_X, train_Y = complete_X[:100], complete_Y[:100]
large_X, large_Y = complete_X[100:], complete_Y[100:]

loss_func = mean_squared_error
```

Use `train_X` and `train_Y` as training samples for ERM. `large_X` and `large_Y` are for the approximation of true data distribution of  $\mathbf{X}$  and  $Y$ , in order to estimate true risk. A plot of a small portion of the dataset.

Use mean squared error as loss function in this problem (where  $\ell(y_i, h(x_i)) = (y_i - h(x_i))^2$ ), unless noted otherwise.

a.(10 points) Let's first define  $\mathcal{H}_k$  to be a collection of all possible polynomial functions of degree  $k$ . Implement the ERM process to select the predictor  $\hat{h} \in \mathcal{H}_k$  with the lowest empirical risk. (Hint: `polyfit` function in numpy could be useful.)



b.(10 points) Experiment with the ERM you built with  $k$  of  $\mathcal{H}_k$  ranging from 0 to 30. Report empirical loss and plot a graph of empirical risk v.s.  $k$ .

c.(10 points) We will further explore the approximation vs. estimation trade-off. First, use the noise-free distribution in data\_generator to estimate  $R^*$  with the **complete dataset**. i.e., the complete dataset has noisy  $(x, y)$  pairs and we know that without noise

$$y^* := \mathbb{E}[Y \mid \mathbf{X} = x] = \cos(0.5 + e^{-x}) + \frac{1}{1 + e^{-x}}. \quad (1)$$

(Note that in real applications, you normally do not have access to the true distribution of  $\mathbf{X}$  and  $Y$ .) Now use `large_X`, `large_Y` to estimate the Bayes risk  $R^*$ , the risk of the ERM for each  $k$ , the estimation error, and the approximation error. Plot graphs of these errors v.s.  $k$ . Experiment with  $k$  range from 0 to 25. (Note: these different errors might not be in the same scale. You can plot one graph for each error v.s.  $k$ )

(a)

```
#python
def evaluate(clf, test_x, test_y, loss_func):
    y_hat = clf(test_x)
    loss = loss_func(test_y, y_hat)
    return loss

def erm(k, train_x, train_y, loss_func):
    clf = np.polyfit(train_x, train_y, k)
    clf = np.poly1d(clf)
    empirical_loss =
        evaluate(clf, train_x, train_y, loss_func)
    return clf, empirical_loss
```

(b)


```
#python
empirical_loss_list = []
loss_func = mean_squared_error
for k in range(30):
```

```

_, emp_loss = erm(k, train_X, train_Y, loss_func)
empirical_loss_list += emp_loss,

plt.clf()
x = range(30)
plt.plot(x, empirical_loss_list)
plt.show()

```



(c)

```

#python
def compute_r_star(x, y, loss_func):
    y_hat = np.cos(0.5 + np.exp(-x)) + 1/(1 + np.exp(-x))
    risk = loss_func(y, y_hat)
    return risk

r_star = compute_r_star(complete_X, complete_Y, loss_func)

emp_risk_list = []
estimation_list = []
approximation_list = []
real_risk_list = []
erm_risk_list = []
max_k = 25

for k in range(max_k):
    clf, empirical_risk = erm(k, train_X, train_Y, loss_func)
    emp_risk_list += empirical_risk,

    erm_risk = evaluate(clf, large_X, large_Y, loss_func)
    erm_risk_list += erm_risk,

    clf, real_risk = erm(k, large_X, large_Y, loss_func)

```

```

estimation_error = erm_risk - real_risk
approximation_error = real_risk - r_star

estimation_list += estimation_error,
approximation_list += approximation_error,
real_risk_list += real_risk,

```

